

## **Interview Q&A Guide: ESG GenAI Data Intern Role**

**Q: What is the difference between a list and a tuple in Python?**

A: Lists are mutable (can be changed), while tuples are immutable (cannot be changed).

**Q: How do you handle missing data in a dataset using Python?**

A: Using pandas functions like `dropna()`, `fillna()`, or interpolation methods.

**Q: What libraries do you use for data analysis?**

A: Common libraries are pandas, NumPy, matplotlib, seaborn, and scikit-learn.

**Q: How would you read and process a CSV file in Python?**

A: Use pandas: `df = pd.read_csv('file.csv')`, then clean or transform as needed.

**Q: Explain the difference between `.loc[]` and `.iloc[]` in pandas.**

A: `.loc[]` accesses rows by labels; `.iloc[]` accesses rows by index position.

**Q: How can you filter rows in a DataFrame based on a condition?**

A: Use a boolean condition, e.g., `df[df['column'] > 100]`.

**Q: How would you merge or join two DataFrames?**

A: Use `pd.merge()` for relational joins, or `pd.concat()` to stack them.

**Q: What is data validation? How would you approach validating data in Python?**

A: Data validation means checking data correctness. Use checks like data types, ranges, or regex for formats.

**Q: How do you handle JSON data in Python?**

A: Use the `json` module to load or dump JSON, or `pandas.read_json()` for tabular formats.

**Q: How do you write a simple Python script to find duplicate values in a list or dataset?**

A: Use `collections.Counter` or `pandas.duplicated()` to identify duplicates.

**Q: What is prompt engineering?**

A: Its the process of designing input prompts to get better, targeted responses from language models.

## **Interview Q&A Guide: ESG GenAI Data Intern Role**

**Q: How does prompt engineering affect the output of an LLM like ChatGPT?**

A: It changes clarity, structure, and relevance of the response. Good prompts give better results.

**Q: Have you ever tried refining prompts to get better results? Can you give an example?**

A: Yes, for summarization tasks I changed 'Summarize this:' to 'Give a 3-bullet summary of:' to improve clarity.

**Q: What is the difference between zero-shot, one-shot, and few-shot prompting?**

A: Zero-shot: no examples; one-shot: 1 example; few-shot: multiple examples to guide the model.

**Q: How would you prompt an LLM to extract structured information from an unstructured paragraph?**

A: Use a clear format like: 'Extract the following fields: Name, Age, Location from this text:'

**Q: What are some challenges you've faced while using AI tools like ChatGPT or OpenAI API?**

A: Challenges include hallucination, inconsistent outputs, and sensitivity to prompt phrasing.

**Q: How can you make a prompt more specific or task-focused?**

A: Be explicit about format, tone, length, and include examples or structure.

**Q: How do you test if a prompt is working well?**

A: Check accuracy, consistency, and repeatability of outputs on varied inputs.

**Q: What tools or platforms have you used to experiment with LLMs (e.g., OpenAI Playground, Hugging Face)?**

A: I've used ChatGPT, OpenAI Playground, and Hugging Face's transformers and inference API.

**Q: Can you give an example of a prompt that failed and how you fixed it?**

A: A vague 'Write about climate change' gave generic results; rephrased to 'List 3 key causes of climate change with 1-line explanation each' improved it.